

Demo day card

One page. Print it or keep it open. Everything else is in `PRESENTATION_GUIDE.pdf`.

Start the system

Two commands, from the project root. Nothing else is needed and nothing reaches the network.

```
cd web && npm run build && cd ..  
NPN MOCK=0 .venv/Scripts/python.exe -m uvicorn server.main:app --port 8000
```

Open <http://127.0.0.1:8000>

Wait for the first prediction to be slow — the model loads on first request, about 15 seconds. **Do one throwaway prediction before the panel walks in** so the load is already paid for.

Check before you present

CHECK	EXPECTED
Rail, bottom left	MAE 5.64, Baseline 11.34
Header	NO amber "SYNTHETIC MODEL" badge
Header	NO red "RESEARCH DEMO" banner
Model evidence tab	charts render, geriatric row in red

If the synthetic badge is showing, you started without `NPN MOCK=0`. Stop and restart.

The run

#	DO	SAY
1	Upload <code>samples/child_08.jpg</code>	True age 8. Reads 9.3, paediatric, rejected for adult trial eligibility. Correct call.
2	Upload <code>samples/teen_17.jpg</code>	True age 17. It reads 24.2 – wrong. But the range is 17.5 to 31.5 and crosses the age-18 line. Watch the ladder light that boundary.
3	Pause	This is the case that matters. The model was wrong and the system refused to act.
4	Review queue	One case waiting, with the rule that routed it.
5	Enter a reviewer ID, override to 17, submit	A human decides. The original model output is kept, not overwritten.
6	Audit trail	Every step. Image fingerprint, never the image.
7	Model evidence	Accuracy, calibration, per-band error with geriatric in red.

Hold on steps 2 and 3 longer than feels comfortable.

Do not use `adult_34.jpg` **or** `senior_72.jpg`. One is sunglasses in a car, the other is mislabelled in the source dataset. Both make the model look worse than it is for reasons that have nothing to do with the model.

If it breaks

Switch to `docs/demo/demo-manual.mp4`. Say **"let me show you the recorded run"** and carry on. Do not narrate the failure, do not apologise. The panel only notices if you make it an event.

If asked to try it on someone's face

Decline. Say: **"we won't run biometric inference on someone who hasn't consented, especially in a healthcare framing."** That answer is a point in your favour.

Numbers, if you blank

MAE	5.64 yr – honest range 5.64 to 6.57
Baseline	11.34 yr (guessing the average age)
Within 5 years	59.3%
Correct band	66.8%
Test set	47,568 held-out images
Best band	adult 4.80
Worst band	geriatric 12.25, and 90+ is 15.95 on 77 images
Confidence works	low-confidence predictions are 3.89x worse, monotonic across all 10 deciles
Interval honesty	an 80% range contains the truth 79.1% of the time
Deferring 15%	error drops 5.64 to 4.99
Fairness	error highest on Black faces in 4 of 5 age bands; worst gap 3.88 yr in 0-17
Routing	15.7% of Black subjects go to review, against 4.1% for the best-served group

Never invent a number. "I'd have to check" is a complete answer.

The one sentence

The model is wrong sometimes. What matters is that it knows when, and stops.